

```
In [17]: from lesson_functions import *
import glob
from os.path import join
from sklearn.metrics import precision_score, classification_report, accuracy_score
from sklearn.svm import LinearSVC, SVC, NuSVC
from sklearn.preprocessing import StandardScaler
import time
import pickle
```

```
In [2]: def train_test_split(X, y, test_size=0.2):
        i = int((1 - test_size) * X.shape[0]) + 1
        X_train, X_test = np.split(X, [i])
        y_train, y_test = np.split(y, [i])
        return X_train, X_test, y_train, y_test
```

## Parameters

While a lot of the parameters were taken from lessons in the class, others came from forum discussions, and others were purely from experimentation.

```
In [3]: color_space = 'HLS'  # Can be RGB, HSV, LUV, HLS, YUV, YCrCb
orient = 9  # HOG orientations
pix_per_cell = 16  # HOG pixels per cell
cell_per_block = 2  # HOG cells per block
hog_channel = "ALL"  # Can be 0, 1, 2, or "ALL"
spatial_size = (16, 16)  # Spatial binning dimensions
hist_bins = 16  # Number of histogram bins
spatial_feat = True  # Spatial features on or off
hist_feat = True  # Histogram features on or off
hog_feat = True  # HOG features on or off
```

## Data Source

```
In [4]: data_dir = "./images/"
cars = glob.glob(join(data_dir, "vehicles/*/*.png"))
notcars = glob.glob(join(data_dir, "non-vehicles/*/*.png"))
```

## Extract Features for Cars and Non-Cars Using Class-Provided Methods and Passing in Appropriate Parameters

```
In [5]: car_features = extract_features(cars, color_space=color_space,
                                       spatial_size=spatial_size, hist_bins=hist_
                                       bins,
                                       orient=orient, pix_per_cell=pix_per_cell
                                       ,
                                       cell_per_block=cell_per_block,
                                       hog_channel=hog_channel, spatial_feat=sp
                                       atial_feat,
                                       hist_feat=hist_feat, hog_feat=hog_feat)
notcar_features = extract_features(notcars, color_space=color_space,
                                   spatial_size=spatial_size, hist_bins=hist_
                                   bins,
                                   orient=orient, pix_per_cell=pix_per_cell
                                   ell,
                                   cell_per_block=cell_per_block,
                                   hog_channel=hog_channel, spatial_feat=
                                   =spatial_feat,
                                   hist_feat=hist_feat, hog_feat=hog_feat)
```

```
/Users/seth.bunke/anaconda/envs/carnd-term1/lib/python3.5/site-package
s/skimage/feature/_hog.py:119: skimage_deprecation: Default value of `b
lock_norm`==`L1` is deprecated and will be changed to `L2-Hys` in v0.15
'be changed to `L2-Hys` in v0.15', skimage_deprecation)
```

## Split Data

```
In [7]: car_X_train, car_X_test, car_y_train, car_y_test = \
        train_test_split(np.array(car_features), np.ones(len(car_features)),
        test_size=0.2)

noncar_X_train, noncar_X_test, noncar_y_train, noncar_y_test = \
        train_test_split(np.array(notcar_features), np.zeros(len(notcar_feat
        ures)), test_size=0.2)

X_train = np.vstack((car_X_train, noncar_X_train))
y_train = np.hstack((car_y_train, noncar_y_train))
X_test = np.vstack((car_X_test, noncar_X_test))
y_test = np.hstack((car_y_test, noncar_y_test))
```

## Scaling the Data

```
In [8]: X_scaler = StandardScaler().fit(np.vstack((X_train, X_test)))
X_train = X_scaler.transform(X_train)
X_test = X_scaler.transform(X_test)
```

## Training the Classifier

Through discussions on the class forum as well as my own testing I selected SVC as the model. The comparison of some models is shown below. The tuning of parameters did not make an appreciable difference in accuracy. Some of that tuning was based on this discussion:

<https://stats.stackexchange.com/questions/31066/what-is-the-influence-of-c-in-svms-with-linear-kernel>

(<https://stats.stackexchange.com/questions/31066/what-is-the-influence-of-c-in-svms-with-linear-kernel>)

```
In [13]: def train_model_report_accuracy(model):
          global X_train
          global y_train
          global X_test
          global y_test
          # Check the training time for the SVC
          t = time.time()
          model.fit(X_train, y_train)
          t2 = time.time()
          print(round(t2 - t, 2), 'Seconds to train model...')
          t = time.time()
          p = model.predict(X_test)
          t2 = time.time()
          # Check the prediction time
          print(round(t2 - t, 4), 'Seconds to predict with model...')
          # Check the score of the SVC
          print('Test accuracy: {:.4f}'.format(accuracy_score(y_test, p)))
          print("Test precision: {:.4f}".format(precision_score(y_test, p)))
          print("")
          print(classification_report(y_test, p, digits=4))
          return model
```

```
In [14]: model = LinearSVC(dual=True, C=0.01)
```

```
In [15]: train_model_report_accuracy(model)
```

```
3.97 Seconds to train model...
0.6391 Seconds to predict with model...
Test accuracy: 0.9772
Test precision: 0.9635
```

	precision	recall	f1-score	support
0.0	0.9914	0.9632	0.9771	1793
1.0	0.9635	0.9915	0.9773	1758
avg / total	0.9776	0.9772	0.9772	3551

```
Out[15]: LinearSVC(C=0.01, class_weight=None, dual=True, fit_intercept=True,
                  intercept_scaling=1, loss='squared_hinge', max_iter=1000,
                  multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
                  verbose=0)
```

```
In [21]: model_1 = SVC()
train_model_report_accuracy(model_1)
```

```
54.37 Seconds to train model...
12.6779 Seconds to predict with model...
Test accuracy: 0.9904
Test precision: 0.9826
```

	precision	recall	f1-score	support
0.0	0.9983	0.9827	0.9904	1793
1.0	0.9826	0.9983	0.9904	1758
avg / total	0.9905	0.9904	0.9904	3551

```
Out[21]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
max_iter=-1, probability=False, random_state=None, shrinking=True,
tol=0.001, verbose=False)
```

```
In [22]: from sklearn import tree
model_2 = tree.DecisionTreeClassifier()
train_model_report_accuracy(model_2)
```

```
76.22 Seconds to train model...
0.05 Seconds to predict with model...
Test accuracy: 0.9285
Test precision: 0.8841
```

	precision	recall	f1-score	support
0.0	0.9831	0.8734	0.9250	1793
1.0	0.8841	0.9846	0.9316	1758
avg / total	0.9340	0.9285	0.9283	3551

```
Out[22]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=N
one,
max_features=None, max_leaf_nodes=None,
min_impurity_decrease=0.0, min_impurity_split=None,
min_samples_leaf=1, min_samples_split=2,
min_weight_fraction_leaf=0.0, presort=False, random_state=N
one,
splitter='best')
```

## Train Using All Data

```
In [23]: X = np.vstack((X_train, X_test))
y = np.hstack((y_train, y_test))

t = time.time()
final_model = SVC()
final_model.fit(X, y)
t2 = time.time()
print(round(t2 - t, 2), 'Seconds to train final model...')
```

74.11 Seconds to train final model...

## Save Model

```
In [24]: with open("final_model.p", "wb") as ofile:
pickle.dump([final_model, X_scaler], ofile)
```

In [ ]: